


# Advanced ML — Lecture 2: Linear Regression

Charalampos Giannakis

## Polynomial curve fitting — motivation



We start with a simple but important example: suppose we have data points sampled from a sinusoidal function with some added noise. Concretely, imagine we generate 10 data points from

$$t = \sin(2\pi x) + \text{disturbance}.$$


This reflects a real-world scenario where observations are not perfect: we have an underlying true function, but our measurements are corrupted by noise. The challenge is: How can we build a model that captures the relationship between  $x$  and  $t$ , despite the disturbances?

A key idea is to approximate the function with a polynomial. Many functions can be expressed as power series. For example, the Taylor expansion of the sine function is

$$\sin(x) = x - \frac{x^3}{3!} + \frac{x^5}{5!} - \frac{x^7}{7!} + \frac{x^9}{9!} + \dots$$

which reminds us that even complex curves can be expressed using polynomials of increasing degree.

## The polynomial model and the error we minimize



In general, we define our polynomial model as

$$y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \cdots + w_Mx^M = \sum_{j=0}^M w_jx^j,$$

where  $M$  is the order of the polynomial and  $w_j$  are the coefficients we need to determine.


To measure how well the model predicts the observed targets  $t_n$ , we use the **sum-of-squares error**:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2.$$

Minimizing  $E(\mathbf{w})$  gives the polynomial curve that best fits the data in the least-squares sense.

## Exploring model capacity with polynomial order

### Order 0 (constant model)



We start with the simplest possible model: a flat line.

It ignores any trend in the data and predicts the same value for every input ( $x$ ).


This is the “high-bias” extreme—very simple, very constrained.

**Model.**  $y(x, \mathbf{w}) = w_0$

**Intuition.** If the data oscillate (like a sine wave), a constant can only match the *Expected Value* level.

Errors remain large because the model cannot bend to follow the signal.

### Order 1 (linear model)




Allow one slope. Now the model can tilt and capture a global upward or downward trend, but still cannot follow curvature.

Bias decreases compared to the constant model, yet the fit is still too rigid for periodic structure.

**Model.**  $y(x, \mathbf{w}) = w_0 + w_1x$

**Intuition.** Good as a baseline when the signal is roughly linear; otherwise, systematic errors remain because the model cannot capture bends.

### Order 3 (cubic model)




With quadratic and cubic terms, the curve can bend and flex.

For our noisy sinusoidal data, a cubic already starts to track the main rise and fall without chasing every noisy fluctuation.

**Model.**  $y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + w_3x^3$

**Intuition.** This is a more balanced regime: enough capacity to capture curvature, but still restrained, so variance is moderate.

### Order 9 (high-degree model)



With many degrees of freedom, the polynomial can pass very close to each training point—including noise.


The red curve wiggles aggressively: it *fits the training set* extremely well but behaves wildly between points and near the boundaries.

**Model.**  $y(x, \mathbf{w}) = w_0 + w_1x + w_2x^2 + \dots + w_9x^9$

**Intuition.** This is the “high-variance” extreme—low training error, high risk of **overfitting**. It motivates the need for **regularization** and for choosing model complexity carefully.

### When models become too flexible: overfitting

#### Training vs. test error



To understand overfitting, we track how error behaves as we increase the polynomial order ( $M$ ).

We measure error using the **root mean square (RMS)**:

$$E_{\text{RMS}} = \sqrt{\frac{2E(\mathbf{w}^*)}{N}}$$

Here,  $E(w^*)$  is the minimized sum-of-squares error and  $N$  is the number of data points.

- On the **training set** (blue curve), the error keeps going down as we add more terms.
- But on the **test set** (red curve), after a certain point the error *increases*.

This means the model has started to memorize noise in the training data rather than learning the true underlying function.

That is the essence of **overfitting**.

### What happens to the coefficients?

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

Looking at the actual values of the fitted coefficients tells the same story.

- For small ( $M$ ), coefficients are small and stable.
- For higher orders like ( $M = 6$ ) or ( $M = 9$ ), the coefficients explode to extremely large positive and negative numbers.

This instability is a direct sign that the model is trying too hard to wiggle through every training point.

It finds a curve that *looks perfect on training data* but generalizes very poorly.


---

**Takeaway.** More complexity is not always better. There's a trade-off:

- Too simple  $\rightarrow$  underfitting (high bias).
- Too complex  $\rightarrow$  overfitting (high variance).
- The sweet spot is somewhere in between, and we'll later see that **regularization** helps us control this balance.

## The effect of dataset size

Small dataset: ( $N = 15$ )




Here we fit a **9th-order polynomial** using only **15 points**.  
The model has a lot of flexibility compared to the limited data available.

- The red curve twists and bends, trying to follow each of the few noisy points.
- As a result, it looks unstable and does not represent the smooth sine wave well.

This is a classic case where too much model complexity meets too little data.

Larger dataset: ( $N = 100$ )



Now we fit the *same* 9th-order polynomial, but with **100 points** instead.  
Notice how different the result is:

- With more data, the curve (red) stays close to the true underlying function (green).
  - The polynomial still has high capacity, but the abundance of data **anchors it** and prevents wild oscillations.
- 

**Takeaway.** The impact of overfitting depends not only on model complexity but also on the **amount of data** available.

- High complexity with few data points → overfitting.
- High complexity with many data points → still manageable, because the data provide enough constraints.

This highlights why in practice, collecting more data can be just as important as designing the right model.

## Regularization: controlling complexity

### Why regularization?

When we used high-degree polynomials, we saw the coefficients  $w_j$  exploding to very large values. This instability is a hallmark of overfitting.

To counter this, we add a penalty that discourages large coefficients.  
The new error function becomes:

$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - y(x_n, \mathbf{w})\}^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2$$


- The first term is the data fitting error.
- The second term is the **regularization penalty**, controlled by  $(\lambda)$ .

Here,  $\lambda$  is a hyperparameter:

- Small  $\lambda \rightarrow$  model behaves like standard least squares.
- Large  $\lambda \rightarrow$  coefficients shrink toward zero.




## Regularization in practice



When  $\ln \lambda = -18$ , the penalty is tiny.

The curve still fits the training data closely and can wiggle quite a lot.

Regularization is present, but not strong enough to really simplify the model.



When  $\ln \lambda = 0$ , the penalty is very strong.

Now the red curve collapses into something almost flat, ignoring the sinusoidal structure completely.

This is the opposite extreme: **underfitting** caused by too much regularization.

---


**Takeaway.** Regularization helps keep coefficients under control and reduces overfitting.

But just like model complexity, the strength of the penalty must be tuned:

- Too little  $\rightarrow$  still overfits.
- Too much  $\rightarrow$  underfits.
- Just right  $\rightarrow$  balances bias and variance.

## Regularization and its effect on errors

$E_{\text{RMS}}$  versus  $\ln \lambda$



Here we plot the **RMS error** against the regularization strength  $\ln \lambda$ .

- For **very small**  $\lambda$  (large negative  $\ln \lambda$ ), the training error (blue) is almost zero, but the test error (red) is high  $\rightarrow$  overfitting.
- As  $\lambda$  increases, the test error goes down, reaching a **sweet spot** where generalization is best.
- If  $\lambda$  becomes too large, both training and test errors rise again  $\rightarrow$  underfitting.

This curve illustrates the classic **bias–variance trade-off**: the right amount of regularization minimizes test error.

### What happens to the coefficients?

	$\ln \lambda = -\infty$	$\ln \lambda = -18$	$\ln \lambda = 0$
$w_0^*$	0.35	0.35	0.13
$w_1^*$	232.37	4.74	-0.05
$w_2^*$	-5321.83	-0.77	-0.06
$w_3^*$	48568.31	-31.97	-0.05
$w_4^*$	-231639.30	-3.89	-0.03
$w_5^*$	640042.26	55.28	-0.02
$w_6^*$	-1061800.52	41.32	-0.01
$w_7^*$	1042400.18	-45.95	-0.00
$w_8^*$	-557682.99	-91.53	0.00
$w_9^*$	125201.43	72.68	0.01

Looking at the fitted coefficients confirms the story:

- With  $\ln \lambda = -\infty$  (no regularization), the coefficients explode to extreme values.
  - With  $\ln \lambda = -18$ , the coefficients are much smaller but still have some variation.
  - With  $\ln \lambda = 0$ , nearly all coefficients shrink to values close to zero.
- 


### Takeaway.

Regularization keeps the model stable by **shrinking coefficients**.

- Without it, high-degree polynomials produce unstable and extreme weights.
- With it, the model is constrained, leading to better generalization.
- The art lies in choosing  $\lambda$  carefully: not too weak, not too strong.

## A deeper analysis

What is the issue?



We began with the sine function and approximated it using polynomials, such as in its Taylor expansion:

$$\sin(x) = x - \frac{x^3}{3!} + \frac{x^5}{5!} - \frac{x^7}{7!} + \frac{x^9}{9!} + \dots$$

This worked for illustrating curve fitting, but it hides a deeper point:

**polynomials are just one choice of functions to build our model.**

In fact, we can frame the whole problem in a much more general way.

## Linear basis function models


The general model is written as:

$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(x) = \mathbf{w}^\top \phi(x)$$

- $\phi_j(x)$  are **basis functions**: known functions we choose beforehand.
- Typically,  $\phi_0(x) = 1$  so that  $w_0$  acts as a bias term.

The word *basis* comes from linear algebra: the functions  $\phi_j(x)$  are like vectors that span a space, and we combine them with weights  $w_j$  to approximate more complex patterns.

## Polynomial basis functions




One familiar choice is **polynomial basis functions**:

$$\phi_j(x) = x^j$$

- These are **global functions**, meaning each  $\phi_j(x)$  affects the output across the entire input space.

- Polynomials can approximate smooth functions well, but they also tend to produce large oscillations when  $M$  is high.

### Gaussian basis functions




Another choice is **Gaussian basis functions**:

$$\phi_j(x) = \exp\left(-\frac{(x - \mu_j)^2}{2s^2}\right)$$

- These are **local functions**: each Gaussian is concentrated around a center  $\mu_j$ .
- The parameter  $s$  controls the width (scale) of the bump.
- With many such bumps, we can flexibly model nonlinear patterns while keeping each basis function localized.

## Sigmoidal basis functions



We can also use **sigmoidal basis functions**:

$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right), \quad \sigma(a) = \frac{1}{1 + \exp(-a)}$$

- These functions transition smoothly from 0 to 1.
- By shifting  $\mu_j$  and scaling with  $s$ , we place these transitions at different points along the input space.
- Sigmoids are widely used in neural networks as activation functions.

---

### Takeaway.


Polynomial fitting was just the beginning.

By introducing **basis functions**, we open the door to a wide range of models.

Depending on whether we choose polynomials, Gaussians, or sigmoids, we can capture very different behaviors—global trends, local bumps, or smooth transitions.

## Maximum likelihood

### Adding noise to our model



So far, we treated our model as deterministic:  $t = y(x, \mathbf{w})$ .  
But in reality, data is noisy. We model this as:

$$t = y(x, \mathbf{w}) + \epsilon$$

where the noise  $\epsilon$  follows a Gaussian distribution:

$$p(\epsilon | \beta) = \mathcal{N}(\epsilon | 0, \beta^{-1})$$

Here,  $\beta$  is the **precision** (inverse variance):  
 $\beta = 1/\sigma^2$ .

---

## Properties of the Gaussian

The Gaussian distribution is defined as:

$$\mathcal{N}(x | \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left( -\frac{(x - \mu)^2}{2\sigma^2} \right)$$

It has mean  $\mu$  and variance  $\sigma^2$ .

Some key properties:

- Always positive:  $\mathcal{N}(x | \mu, \sigma^2) > 0$

- Normalized:  $\int_{-\infty}^{\infty} \mathcal{N}(x | \mu, \sigma^2) dx = 1$

- Expectations:

$$\begin{aligned} E[x] &= \mu, \\ E[x^2] &= \mu^2 + \sigma^2, \\ \text{var}[x] &= \sigma^2. \end{aligned}$$

Proofs are in Lecture 1 (pp. 3–4)

## Likelihood of data points

Assume observations from a deterministic function with added Gaussian noise:

$$t = y(x, \mathbf{w}) + \epsilon \text{ where } p(\epsilon | \beta) = \mathcal{N}(\epsilon | 0, \beta^{-1})$$

This is the same as saying:

$$p(t | x, \mathbf{w}, \beta) = \mathcal{N}(t | y(x, \mathbf{w}), \beta^{-1})$$

So each observed target  $t$  is normally distributed around our model prediction  $y(x, \mathbf{w})$  with variance  $\beta^{-1}$ .

Recall:

$$y(x, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(x) = \mathbf{w}^\top \phi(x)$$

This is the same as saying  $p(t | \mathbf{x}, \mathbf{w}, \beta) = \mathcal{N}(t | y(\mathbf{x}, \mathbf{w}), \beta^{-1})$

Given all observed inputs  $\mathbf{X} = \{x_1, \dots, x_N\}$   
and targets  $\mathbf{t} = [t_1, \dots, t_N]^\top$ ,

the likelihood is:

$$p(\mathbf{t} | \mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N \mathcal{N}(t_n | \mathbf{w}^\top \phi(x_n), \beta^{-1})$$

Taking the logarithm gives:

$$\ln p(\mathbf{t} | \mathbf{w}, \beta) = \sum_{n=1}^N \ln \mathcal{N}(t_n | \mathbf{w}^\top \phi(x_n), \beta^{-1})$$

which simplifies to:

$$\ln p(\mathbf{t} | \mathbf{w}, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w})$$



where

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^\top \phi(x_n)\}^2$$

---

### Why squared error?

The error function we minimize is exactly the **sum of squared errors**:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^\top \phi(x_n))^2$$

This arises naturally from the Gaussian noise assumption.

- If errors are Gaussian  $\rightarrow$  squared error is the right measure.
  - If errors follow another distribution (e.g., in finance, heavy-tailed), we would use a different error measure.
- 

### Takeaway.

By assuming Gaussian noise, maximum likelihood estimation for linear basis function models reduces to **minimizing squared error**.

This connects probability theory with the familiar least-squares approach.

## Maximum Likelihood and Regularization

When we try to optimize the weights  $\mathbf{w}$  using maximum likelihood, we take the gradient of the log-likelihood and set it to zero:

$$\nabla_{\mathbf{w}} \ln p(t|\mathbf{w}, \beta) = \beta \sum_{n=1}^N \{t_n - \mathbf{w}^\top \phi(\mathbf{x}_n)\} \phi(\mathbf{x}_n)^\top = 0$$

Setting the gradient to zero gives the stationary condition:  $\sum_{n=1}^N \phi(x_n) (t_n - \mathbf{w}^\top \phi(x_n)) = 0$ .

Drop the positive constant  $\beta$  and expand:

$$\sum_{n=1}^N t_n \phi(\mathbf{x}_n) - \sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^\top \mathbf{w} = \mathbf{0}.$$

Rearrange:

$$\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^\top \mathbf{w} = \sum_{n=1}^N t_n \phi(\mathbf{x}_n).$$

**Matrix notation.**

Let the design matrix  $\Phi \in \mathbb{R}^{N \times M}$  have rows  $\phi(\mathbf{x}_n)^\top$  and let  $\mathbf{t} = (t_1, \dots, t_N)^\top$ . Then

$$\sum_{n=1}^N \phi(\mathbf{x}_n) \phi(\mathbf{x}_n)^\top = \Phi^\top \Phi, \quad \sum_{n=1}^N t_n \phi(\mathbf{x}_n) = \Phi^\top \mathbf{t}.$$

So we obtain the normal equations:

$$\Phi^\top \Phi \mathbf{w} = \Phi^\top \mathbf{t}.$$

Here  $\Phi \in \mathbb{R}^{N \times M}$  has rows  $\phi(\mathbf{x}_n)^\top$ ,  $\mathbf{t} \in \mathbb{R}^N$ , and  $\mathbf{w} \in \mathbb{R}^M$ .

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \cdots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \cdots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \cdots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

**Closed-form solution (when invertible)**

$$\mathbf{w}_{\text{ML}} = (\Phi^\top \Phi)^{-1} \Phi^\top \mathbf{t}$$

**If not invertible (or ill-conditioned)**

Use the Moore–Penrose pseudoinverse:  $\mathbf{w} = \Phi^+ \mathbf{t}$ .

(When  $\Phi$  has full column rank,  $\Phi^+ = (\Phi^\top \Phi)^{-1} \Phi^\top$ .)


**Ridge regularization**

$$(\Phi^\top \Phi + \lambda I) \mathbf{w} = \Phi^\top \mathbf{t} \Rightarrow \mathbf{w}_{\text{ridge}} = (\Phi^\top \Phi + \lambda I)^{-1} \Phi^\top \mathbf{t}.$$

**Notes**

- Start from the normal equations:  $\Phi^\top \Phi \mathbf{w} = \Phi^\top \mathbf{t}$ .
- $\Phi^\top \Phi$  is symmetric and positive semidefinite.
- It is invertible **iff** the columns of  $\Phi$  are linearly independent (full column rank).
- $\Phi$  is the design matrix of basis functions evaluated at the data points.

## Geometric view



- Building the **design matrix**  $\Phi$ , each column (basis vector) is  $\phi_j = (\phi_j(\mathbf{x}_1), \dots, \phi_j(\mathbf{x}_N))^T \in \mathbb{R}^N$ .
- Let  $\mathcal{S} = \text{span}\{\phi_1, \dots, \phi_M\} \subseteq \mathbb{R}^N$  (the **column space** of  $\Phi$ ).
- For any weights  $\mathbf{w}$ , the model outputs on the training data are  $\mathbf{y} = \Phi \mathbf{w}$ , hence  $\mathbf{y} \in \mathcal{S}$ .
- **Least squares** picks  $\mathbf{w}_{\text{ML}}$  so that  $\mathbf{y} = \Phi \mathbf{w}_{\text{ML}}$  is the **orthogonal projection** of  $\mathbf{t}$  onto  $\mathcal{S}$ :

$$\mathbf{y} = \arg \min_{\mathbf{z} \in \mathcal{S}} \|\mathbf{t} - \mathbf{z}\|^2.$$

- The residual  $\mathbf{r} = \mathbf{t} - \mathbf{y}$  is orthogonal to every column of  $\Phi$ :

$$\Phi^T \mathbf{r} = \mathbf{0} \iff \Phi^T \Phi \mathbf{w}_{\text{ML}} = \Phi^T \mathbf{t}.$$

- **Geometric picture:**  $\mathcal{S}$  is the subspace spanned by the basis columns (a plane).  $\mathbf{y}$  is the foot of the perpendicular from  $\mathbf{t}$  to  $\mathcal{S}$  (closest point).  
Pythagoras:  $\|\mathbf{t}\|^2 = \|\mathbf{y}\|^2 + \|\mathbf{r}\|^2$ .

**One-liner:** least squares projects  $\mathbf{t}$  onto  $\text{col}(\Phi)$ ; the projection is  $\mathbf{y} = \Phi \mathbf{w}_{\text{ML}}$ .

## Regularization

### Why regularize?

Our plain least-squares fit can overfit when  $M$  is large, features are collinear, or  $N$  is small. We keep the data fit but **discourage large weights** by adding a penalty:

$$E(\mathbf{w}) = E_D(\mathbf{w}) + \lambda E_W(\mathbf{w}).$$

With squared-error data term and a quadratic penalty:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^\top \phi(\mathbf{x}_n))^2 + \frac{\lambda}{2} \mathbf{w}^\top \mathbf{w}.$$


Minimizing gives the **ridge** solution

$$\mathbf{w}_{\text{ridge}} = (\lambda I + \Phi^\top \Phi)^{-1} \Phi^\top \mathbf{t}.$$

- $\lambda \uparrow \rightarrow$  stronger shrinkage, lower variance, higher bias.
- $\lambda \downarrow \rightarrow$  solution moves back toward least squares.

---

### A more general view




Use an  $L_q$  penalty:

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^\top \phi(\mathbf{x}_n))^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q.$$

- $q = 2$  (quadratic)  $\rightarrow$  **ridge**. Constraint sets are circles (in 2D).
- $q = 1$  (absolute value)  $\rightarrow$  **lasso**. Constraint sets are diamonds.
- $q < 1$  makes even pointier shapes (nonconvex).

---

## Geometry & sparsity



Think of **data error contours** (blue ellipses) and a **regularization ball** (green/red boundary).

- Ridge ( $q = 2$ ): the circular boundary usually touches an ellipse away from the axes, so both coordinates are nonzero. We get shrinkage but not sparsity.
- Lasso ( $q = 1$ ): the diamond has sharp corners on the axes. The first point of contact is often a corner, so one or more  $w_j$  become exactly 0.  
Result: sparser solutions (feature selection).

**Takeaway:** regularization balances fit and simplicity. Ridge stabilizes and shrinks; lasso often sets coefficients exactly to zero.